USING SENTIMENT SCORES FOR STANCE DETECTION ON CLIMATE-CHANGE STATEMENTS

Michael Zhu
UC Berkeley School of Information
michael.zhu@berkeley.edu
Spring 2021

ABSTRACT

Despite the catastrophic effects that climate change is forecasted to have over the next century, the topic still remains a partisan issue in America. The ability to automatically and accurately determine the opinion of a piece of text about climate change could be useful for numerous applications, such as tracking public opinion or identifying actors that consistently reject or downplay the validity of scientific evidence. Stance detection is a popular classification task in natural language processing that is well suited for developing models to determine the opinion of a piece of text. In this study, I sought to improve stance detection performance on a recently curated data set (DeSMOG), containing 2,050 global-warming-related sentences from various news sources. Specifically, I focused on engineering new features and tested the hypothesis that adding sentiment scores as a feature to BERT embeddings would improve classification accuracy and macro-F1 score. Using two different methods for sentiment scoring, I find that neither were able to reliably improve performance above a baseline model, which used only the pooled output embedding from BERT (0.70 accuracy and 0.71 macro-F1). I discuss possible reasons for this null result as well as broader connections between stance detection and sentiment analysis.

1 Introduction

Support for climate change policy is becoming an increasingly partisan issue in the United States over recent years [1]. Accordingly, many studies have investigated the effects of using different methods to communicate climate-change-related information to the public [2, 3, 4]. A recent study that investigated differences in word choice between various news sources developed a supervised model to perform stance detection (Luo et al., 2020). The authors used the model to label sentences from news articles as *agree*, *disagree*, or *neutral* with respect to a single target statement: "Climate change/global warming is a serious concern." The authors find that a BERT model performed best at this classification task with 0.75 accuracy and 0.73 macro-F1 score. In this project, I aimed to improve upon the performance of BERT by using sentiment scores as an additional feature to perform stance detection on the same data set as in Luo et al., 2020.

2 Background and Methods

Stance detection is a natural language processing (NLP) task that aims to determine the opinion, or "stance," of a piece of text. Specifically, stance detection is most commonly performed by classifying the relationship between a candidate piece of text and a target statement as falling into one of three classes: *agree, disagree, neutral* (see Table 1 for examples). This task has received increased attention in recent years as concerns grow over the spread of misinformation, popularly known as "fake news" [6]. Therefore, much of the recent literature on stance detection has aimed to identify misinformation in text collected from news sources, including conventional news websites as well as social media platforms. In particular, two of the most popular stance detection competitions were the SemEval-2016 Task 6 [7], composed of tweets, and the 2017 Fake News Challenge (FNC) Stage 1 [8], composed of news articles. Notably, in contrast to these data sets that use numerous pairs of candidate and target statements, all examples in the data set used in this study were labeled relative to a single target statement: "Climate change/global warming is a serious concern."

In addition to its potential for detecting misinformation, an effective stance detection model can have various other applications for downstream analyses, such as tracking public opinion on a specific topic or researching how supporters of different sides of an argument may present information differently. An example of the latter: Luo et al., 2020

developed and used a stance detection model to label stances of over 500,000 sentences related to global warming scraped from news articles, and then studied differences in word choice between "global-warming- accepting media" and "global-warming-skeptic media" ; among other findings, they discovered that global-warming-skeptic sources tend to more often use words that cast "opponent-doubt," "(eg. pretend, claim; inaccurate, alleged)." Critically, the authors were only able to perform the primary analyses of the study because of the automatic labeling made possible by the stance detection model. Therefore, just as with detecting misinformation, many future applications of stance detection models would benefit greatly from the model itself being as accurate and precise as possible. So the primary focus of this study is to improve the classification performance (measured by accuracy and macro-F1 score) of the model developed by Luo et al., 2020 by training on the same data set released by the authors.

In particular, I sought to improve stance detection performance by testing the hypothesis that authors who are more opinionated on climate change also write with stronger sentiment. If this were true, then perhaps the sentiment of a piece of text would add useful information for a model to distinguish between different stances. Therefore, a major portion of my analysis focused on deriving sentiment scores for text data and using the scores as features to improve stance detection performance. Using sentiment-related information to improve stance detection has been tested before, but not specifically with climate-change-related text. Many previous studies used the SemEval-2016 Twitter data set and tried a variety of approaches to incorporate sentiment information, including using sentiment lexicons to derive features for an SVM classifier [9], building a "joint neural network model" [10], and using a multi-task learning architecture that performs both sentiment classification and stance detection [11]. Overall, these studies report that incorporating sentiment information improved performance beyond state-of-the-art (SOTA) models for the SemEval-2016 data set. In this study, I use sentiment scores as features for stance detection on the DeSMOG data set by Luo et al., 2020.

For deciding which general model architecture to use, the stance-detection literature seems to clearly point to transformers, specifically BERT, as performing best. Although historically the most popular models include SVMs and RNNs [8, 12, and see 13 for review], a recent study focused on model comparison for stance detection showed that BERT outperformed previous SOTA models [14]. Importantly, the authors of the DeSMOG data set also reported that a BERT-base model performed best at stance detection with 0.75 accuracy and 0.73 macro-F1 score. Therefore, in this study I compared different variants of BERT-based models with and without sentiment features.

Specifically, I make primarily two contributions to previous work:

1. Experimenting with 3 new features used for stance detection on the DeSMOG data set
   (1) Number of tokens
   (2) swnet-sent - sentiment scores derived from using SentiWordNet
   (3) bert-sent - sentiment scores derived from BERT 'sentiment-analysis' pipeline
2. Detecting statistical significance of differences in these features between classes

3 Results

3.1 Data

The data used for all analyses were obtained from the public github repository of Luo et al., 2020[1]. The data set contains 2,050 global-warming-related sentences from various news websites. For modeling, 200 sentences were used as a held-out test set (stratified by label and political leaning of the news source), and the remaining 1,850 sentences were used for training. Some example sentences and corresponding labels are displayed in Table 1.

3.2 Exploratory Data Analysis

The data showed a relatively large class imbalance with disagree stances occurring around half as often as the other two classes {*agree*: 0.37, *disagree*: 0.20, *neutral*: 0.43}. According to previous work on this data set, resampling techniques tended to perform poorly compared to using class weights [5]; therefore, I also implemented class weights to adjust the loss function in all models.

Next, I looked at the distribution of token lengths for the training set to get an idea of how much text is in the sentences and if there are substantial differences between classes. The mean token length across all training examples was 22.4 ± 11.1 (std), and sentences with an *agree* stance were significantly longer than those with a *disagree* stance (Wilcoxon rank-sum, $p = 0.0035$). In addition to overall length of sentences, I also examined the occurrences of specific tokens and how their frequencies differed between classes (Appendix A).

[1] `https://github.com/yiweiluo/GWStance`

Table 1: Target statement and 5 example sentences with labels from the DeSMOG data set.

Target Statement: "Climate change/global warming is a serious concern."	
Example Sentence	**Label**
If carbon dioxide emissions continue to rise beyond 2020, or even remain level, the temperature goals set in Paris become almost unattainable.	neutral
The study is one more example that you can get any answer you want when the thermometer data errors are larger than the global warming signal you are looking for.	disagree
Millions more people around the world are threatened by river floods in coming decades due to climate change.	agree
45% of the general public view perceived global warming as caused by humans	neutral
Two billion people may be displaced by rising sea levels by the turn of the next century.	agree

Overall, from these exploratory analyses, I decided to implement the following for stance detection classification: (1) use class weights in the loss function and (2) test number of tokens ("numTokens") as an additional feature due to its significant difference between classes.

3.3 Baseline Stance Detection Model

As a baseline model, I used a neural network with the following architecture (in order of first hidden layer to output): (1) pre-trained BERT-base-uncased model from the Huggingface Transformers library; (2) a single fully-connected dense layer with 256 neurons; (3) a softmax output layer with 3 units (i.e. the 3 classes: *agree, disagree, neutral*). The model was trained using 5-fold cross validation on 1,850 sentences and achieved validation performance of 0.67 ± 0.02 accuracy (mean ± std across folds) and 0.64 ± 0.03 macro-F1. These scores did not match the best performing model from Luo et al., 2020 (0.75 accuracy and 0.73 macro-F1); however, they did come relatively close and outperformed all other linear models implemented by Luo et al, 2020, thus reaffirming the remarkable capabilities of BERT for transfer learning.

The most apparent issue with the baseline model is overfitting to the training data with training accuracy approaching 1 and validation accuracy increasing in initial epochs but then peaking at ~0.65 accuracy usually by the fourth epoch (Fig. 1). To combat this overfitting, I ran a series of experiments including the following: varying the dropout rate of the dense layer, varying the number of neurons in the dense layer, adding an additional dense layer with dropout, lowering the number of epochs, and changing the proportions of the train/dev split. However, none of these methods were very effective at increasing validation accuracy, and the main purpose of this study is to test the use of sentiment information for stance detection. Therefore, I proceeded to use the same hyperparameters as in the baseline model, acknowledging that overfitting due to model architecture and hyperparameter values may be a confound for interpreting results.

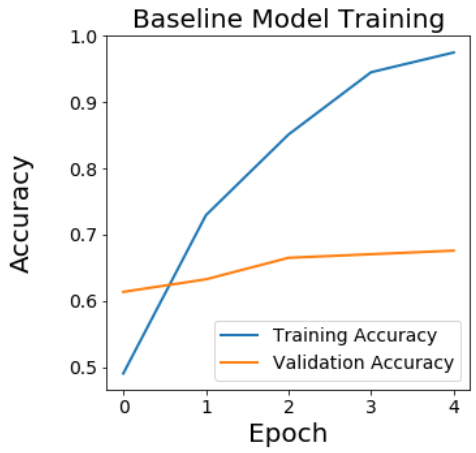

Figure 1: Training and validation accuracy of the baseline model using BERT-base-uncased across 5 epochs (epoch 0 is after finishing first epoch).

3.4 Engineering Sentiment Features

To test the main hypothesis that sentiment scores can be used to improve the performance of BERT for stance detection, I engineered two features using different methods that both aim to capture sentiment information.

The first method uses a lexicon-based approach by using SentiWordNet (swnet) to assign sentiment scores. Swnet comprises a list of annotations for each WordNet synset according to the degree of positivity, negativity, and neutrality (each one on a scale of 0-1 and all three scores sum to 1). Swnet was implemented using the nltk library. For each example, a single scalar value was calculated (see Appendix Methods for details) to capture the overall sentiment of the sentence (Fig. 2, *Top Row*).

The second sentiment feature was derived from the probabilities outputted from the "sentiment-analysis" pipeline of the Transformers library. This implementation uses a pre-trained distilBERT model to classify the sentiment of a piece of text as one of two classes: positive, negative and also outputs the probability predicted for each class. These probabilities were calculated for each example sentence and then transformed (see Appendix Methods) to use as a second sentiment feature for stance detection classification (Fig. 2, *Bottom Row*).

Though the shapes of the distributions for these two sentiment features are starkly different, both features show that *disagree* sentences tend to have lower sentiment scores than *agree* sentences. These differences are statistically significant for both features (for swnet-sent: unpaired t-test, $p = 0.011$, for bert-sent: Wilcoxon rank-sum, $p < 1\text{x}10^{-9}$), though the effect sizes are only small-moderate: for swnet-sent, Cohen's D = 0.154, for bert-sent, Cohen's D = 0.308.

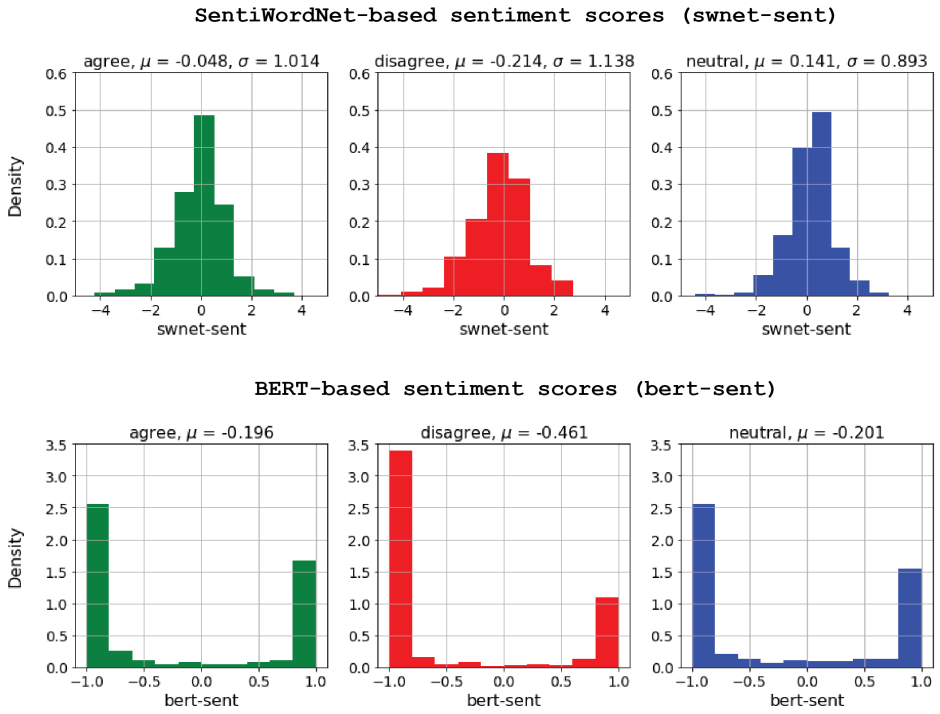

Figure 2: Distributions of the two sentiment features, swnet-sent (*top*) and bert-sent (*bottom*), separated by class.

3.5 Adding New Features to Baseline Model

In total, 3 features were engineered: (1) number of tokens, (2) swnet-sent, and (3) bert-sent. The latter two were the main variables of interest while "number of tokens" was mainly used for exploratory purposes and to control for adding a scalar value onto the BERT embedding. The features were implemented by using the same architecture and hyperparameters as in the baseline model; the only change was concatenating the new feature(s) onto the 768-dimensional pooled output embedding from the pre-trained BERT (eg. if a single feature is being tested, a 769-dimensional vector is inputted into the 256-neuron dense layer and then passed to the softmax output layer). In total, I ran 4 experiments: one for each engineered feature individually and the fourth used all three engineered features together. Table 2 shows the results of all models.

Overall, all models performed within a relatively narrow range of each other and were far better than majority-class prediction. The model using only bert-sent performed best on validation data by a small margin with 0.69 ± 0.02 (std) accuracy and 0.66 ± 0.03 macro-F1. However, qualitatively, this performance did not seem reliably better than the baseline model. Running more cross-validation training sessions would be necessary to determine whether the

difference in evaluation metrics is statistically significant. Evaluating predictions on the held-out test set using the best performing model achieved 0.71 accuracy and 0.70 macro-F1, which falls short of the 0.75 accuracy and 0.73 macro-F1 previously achieved by Luo et al., 2020 (who did not use any explicit sentiment features). Based on these results, using sentiment features did not have a large impact on accuracy or macro-f1 for stance detection on the DeSMOG data set.

Table 2: Summary of model performance: (mean ± std) across 5-fold cross-validation for each model

Model	**Validation Accuracy**	**Validation Macro-F1**
Majority Class	0.43	0.17
BERT (baseline)	0.67 ± 0.02	0.64 ± 0.03
BERT With Only "numTokens"	0.64 ± 0.03	0.65 ± 0.03
BERT With Only swnet-sent	0.69 ± 0.02	0.66 ± 0.03
BERT With Only bert-sent	0.67 ± 0.02	0.65 ± 0.03
BERT With 3 Engineered Features	0.67 ± 0.03	0.64 ± 0.04

3.6 Error Analysis

To better understand the model's classification patterns and provide insight for future analyses, I examined the confusion matrix of the best performing model evaluated on the held-out test set (Fig. 3). By far the most common misclassification (11% of all test examples) was predicting *neutral* when the true label was *agree*. This result corroborates findings reported in Luo et al., 2020, who found that F1 score for the *agree* class was the lowest out of the three classes (i.e. likely due to more false negatives). Qualitatively, from reading some of the misclassified examples it seems like the model doesn't effectively learn relatively abstract relationships between climate change and related entities. For example, "'We can expect the Arctic to be ice-free in summer within 20 years" and "Coal would have to be phased out even before the Paris Agreement to combat climate change." were both labeled *neutral*. In these examples, relationships like climate change and ice melting or coal usage - as well as why they may be a cause for concern - are not obvious and require a certain degree of previous knowledge and abstraction. Therefore, (as with most NLP tasks) future models would likely perform much better with methods that more effectively learn abstract relationships between ideas represented by text.

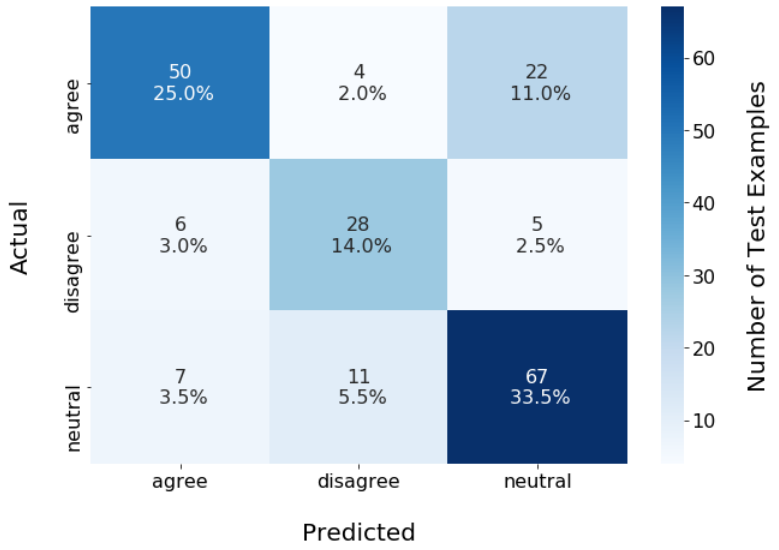

Figure 3: Confusion matrix of predictions on the held-out test set (using the model that performed best based on mean 5-fold cross-validation accuracy).

4 Discussion and Conclusion

My initial hypothesis was that authors who are more opinionated on climate change also write with stronger sentiment. Overall, I find mixed results: opinions that are labeled *disagree* do have significantly lower sentiment scores than opinions labeled *agree*, however, using sentiment scores as features for classification was not able to improve performance on stance detection. Furthermore, there are two important limitations of this study: (1) models tend to overfit to the training data and (2) the engineered sentiment features use relatively crude methods to capture sentiment information. In previous work, methods that focused on changing network architecture to capture sentiment as opposed to engineering single scalar values had more success with improving stance detection on the SemEval2016 data set [10, 11]. Therefore, future analyses may also be able to expand on the work here and still develop more effective methods to use sentiment information. Intuitively, part of the reason for the ineffectiveness of sentiment scores in this study could be due to the relatively large input space of the BERT output embedding. Simply concatenating a scalar sentiment score onto a 768-dimensional embedding may make any information contained in the sentiment score too sparse for the model to extract useful information for separating between classes. This problem would also give more reason for future work to focus on architecture engineering rather than feature engineering, such as perhaps by making modifications to transformers so that the embeddings themselves more explicitly capture sentiment information.

Developing more reliable and accurate stance detection models could have wide ranging applications, especially for automatically tracking public opinion on controversial or important issues such as climate change. For instance, in a previous study that in part inspired this project, Dahal et al., 2019 measured sentiment in over 500 million tweets related to climate change around the globe from 2016 to 2018. Though confounded by using only Twitter data, the study presents time courses of average sentiment that reveal interesting findings, such as a negative spike in sentiment (reaching its lowest point within the 2 year span) following the US withdrawal from the Paris Agreement, and the US consistently having negative sentiment over the 2 year span compared to other countries such as Australia, which had the most consistently positive sentiment. Building on this study, stance detection can play a complementary role by tracking public stance over time in addition to sentiment. Although, as NLP tasks, stance detection and sentiment analysis are often thought of as similar, each one aims to capture a distinct construct and outputs of the two tasks do not always align [see 13 and 15 for discussion]. For example, in the DeSMOG data set used in this study, the sentence "Extreme storms will be much more frequent as a result of global warming" has a strongly negative sentiment but agrees with the stance that climate change is a serious concern. Therefore, future work on tracking climate-change opinions - as well as countless other opinion mining applications - could benefit from also using stance detection to capture emotional valence *relative* to a target opinion as opposed to relying mostly on *overall* emotional polarity captured by traditional sentiment analysis. In this study, I tested if incorporating sentiment information into a stance detection model could improve model performance, but future work could also focus on applying both tasks in tandem.

References

[1] Pew Research Center. (2020). As economic concerns recede, environmental protection rises on the public's policy agenda.

[2] Bolsen, T., Shapiro, M. A., Bolsen, T., Shapiro, M. A. (2017). Strategic Framing and Persuasive Messaging to Influence Climate Change Perceptions and Decisions. In Oxford Research Encyclopedia of Climate Science. Oxford University Press.

[3] 3. Bolsen, T., Palm, R., Kingsland, J. T. (2019). The Impact of Message Source on the Effectiveness of Communications About Climate Change. Science Communication, 41(4), 464–487.

[4] Palm, R., Bolsen, T., Kingsland, J. T. (2020). '"Don't Tell Me What to Do"': Resistance to Climate Change Messages Suggesting Behavior Changes. Weather, Climate, and Society, 12(4), 827–835.

[5] Luo, Y., Card, D., Jurafsky, D. (2020). DeSMOG: Detecting Stance in Media On Global Warming. arXiv preprint arXiv:2010.15149, 2020.

[6] Lazer, D. M. J., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Pennycook, G., Rothschild, D., Schudson, M., Sloman, S. A., Sunstein, C. R., Thorson, E. A., Watts, D. J., Zittrain, J. L. (2018). The science of fake news: Addressing fake news requires a multidisciplinary effort. Science, 359(6380), 1094–1096.

[7] Mohammad, S. M., Kiritchenko, S., Sobhani, P., Zhu, X., Cherry, C. (2016). SemEval-2016 Task 6: Detecting Stance in Tweets.

[8] Hanselowski, A., PVS, A., Schiller, B., Caspelherr, F., Chaudhuri, D., Meyer, C. M., Gurevych, I. (2018). A Retrospective Analysis of the Fake News Challenge Stance Detection Task.

[9] Sobhani, P., Mohammad, S., Kiritchenko, S. (2016). Detecting stance in tweets and analyzing its interaction with sentiment. In Proceedings of the fifth joint conference on lexical and computational semantics, pp. 159–169.

[10] Sun, Q., Wang, Z., Li, S., Zhu, Q., Zhou, G. (2019). Stance detection via sentiment information and neural network model. Frontiers of Computer Science, 13(1), 127–138.

[11] Li. Y, and Caragea C. (2019). Multi-Task Stance Detection with Sentiment and Stance Lexicons. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), pp. 6299–6305, 2019

[12] Rakholia, N., Bhargava, S. (2017). "is it true?" - deep learning for stance detection in news. tech. rep. Stanford Uiversity.

[13] Küçük, D., Fazli, C. A. N. (2020). Stance detection: A survey. In ACM Computing Surveys (Vol. 53, Issue 1). Association for Computing Machinery.

[14] Ghosh, S., Singhania, P., Singh, S., Rudra, K., Ghosh, S. (2020). Stance Detection in Web and Social Media: A Comparative Study.

[15] Dahal, B., Kumar, S. A. P., Li, Z. (2019). Topic modeling and sentiment analysis of global climate change tweets. Social Network Analysis and Mining, 9(1).

5 Appendix A

Fig. S1 shows the tokens that had the largest differences in occurring between *agree* and *disagree* sentences. The token that was most commonly used more in *agree* sentences than *disagree* sentences was a comma (used ~20% more often), which makes intuitive sense based on the previous analyses that identified that *agree* sentences tend to be longer than *disagree* sentences; the use of more commas may also suggest the use of more complicated sentence structures with multiple clauses. Although some other tokens in the list were interesting (eg. *disagree* uses "global" and "warming" more often and *agree* uses "fossil," "sea," "rise," and "we" more often), I did not explicitly use any of these insights for other analyses in this study due to time constraints. Nonetheless, future work may benefit from applying this analysis as well as its extensions, such as using an n-gram approach as opposed to only using unigrams (i.e. single tokens).

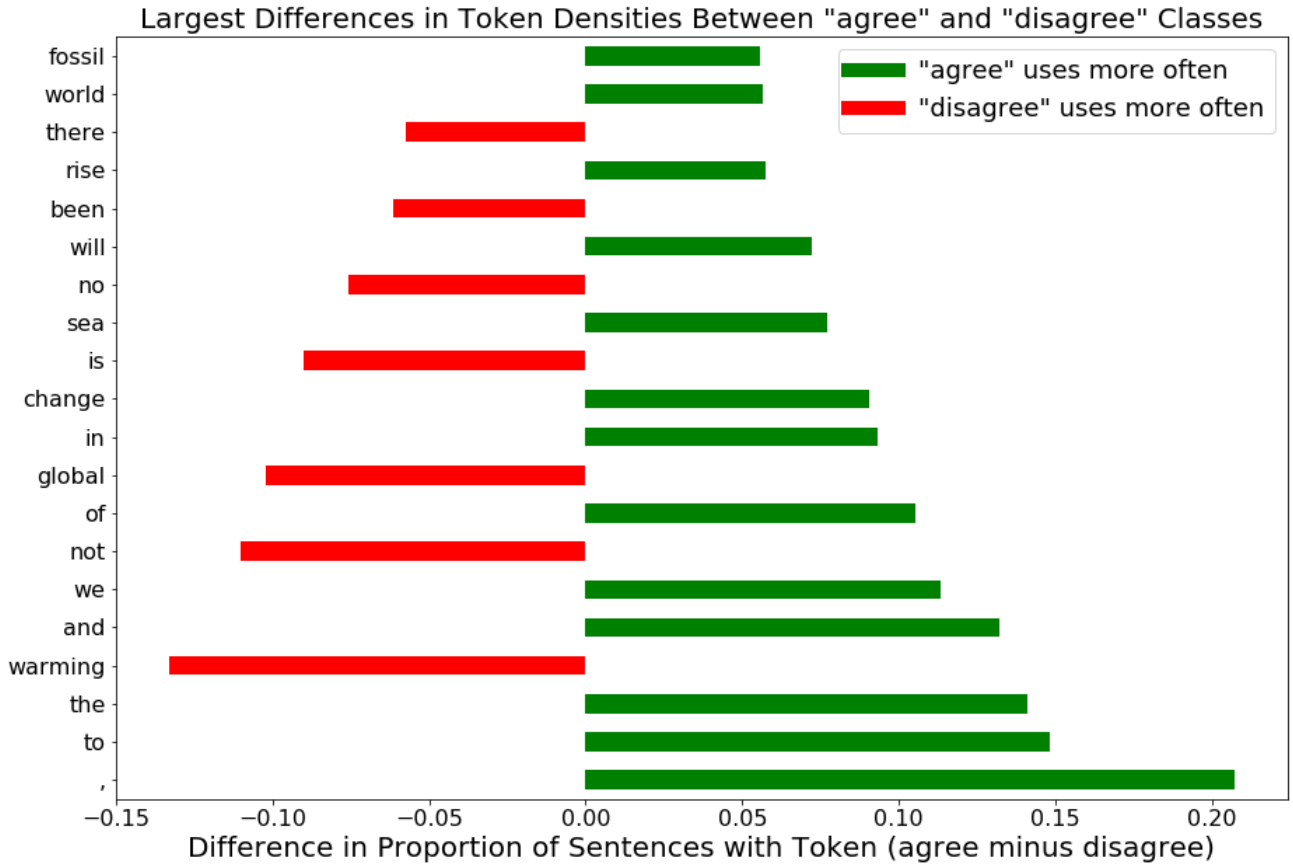

Figure S1: Tokens with the largest differences in frequency used between *agree* and *disagree* examples. Proportions are first calculated within-class and then subtracted.

6 Appendix Methods

Data

All data were obtained from the public github repository of Luo et al., 2020 (https://github.com/yiweiluo/DeSMOG). The data set contains 2,050 global-warming-related sentences scraped from news articles from both "left-leaning" and "right-leaning" media. The most common news sources scraped from include The New York Times, Vox, Mother Jones, Washington Post, The Nation, Guardian (US), Newsmax, Fox, Redstate, Washington Examiner, Washington Times, and Breitbart. Class labels were determined based on annotations made by Amazon Mechanical Turk workers. See Luo et al., 2020 for more details on how the data set was created.

Sentiment Features

SentiWordNet-based scores (swnet-sent) - swnet-sent values were calculated using the following procedure: for each example in the data set, (1) look up the first listed positive and negative SentiWordNet score of each token (tokens not found in the SentiWordNet lexicon were ignored); (2) subtract the negative score from the positive score for each token; (3) calculate the mean of the resulting scores across all tokens, which results in a single scalar sentiment score for each example. All scores across all examples were z-scored to scale the values so that they are comparable to the embedding output of BERT. A key limitation of this method currently is using the first listed scores available from SentiWordNet, which does not take into account the part of speech of each token in the specific example. Due to time constraints, I was not able to implement a part of speech tagger in order to assign more accurate SentiWordNet scores, but future work can likely improve the validity of swnet-sent scores by first pos-tagging all tokens for each example sentence individually.

BERT-based scores (bert-sent) - The "sentiment-analysis" pipeline from the Transformers library was used with default settings. Values for the bert-sent feature were calculated using the following procedure: for each example in the data set, (1) use the pipeline to output two probabilities, one for each of the classes: positive and negative; (2) to capture a "net sentiment" subtract the probability negative from probability positive (so when the model is more uncertain, the value is closer to 0); (3) change the sign of all examples labeled negative by the pipeline to have negative values (i.e. to capture "negative" sentiment).

Modeling

All models used BERT-base-uncased and the following hyperparameters: {BERT max token length: 90, learning rate: 2e-5, number of epochs: 7, batch size: 16, number of neurons in dense layer: 256, dropout rate in dense layer: 0.1}. Accuracy and Macro-F1 were chosen as the primary evaluation metrics due to the class imbalance as well as based on previous work using this data set [5].

All models were trained and evaluated on a NVIDIA GeForce RTX 2060 (6GB VRAM) using the NVIDIA CUDA interface.